

Auteur Skill

> Ported from [agiwhitelist/auteur](https://github.com/agiwhitelist/auteur) (MIT), snapshot

> commit [9bca227d]

(https://github.com/agiwhitelist/auteur/commit/9bca227df9877e60dc45d49783c8cbd885eccd9b)

> — see `LICENSE` . Scripts, templates and references are the upstream files (CRLF→LF), with

> Hermes adaptation notes and `references/` path fixes as the only edits.

Auteur designs and builds web experiences the way a film director makes a film: script first, then assets, then the shoot, then the cut. It has three registers — **build** (an excellent conventional site), **direct** (a cinematic scroll-directed site) and **system** (a multi-screen product as one design system) — on one shared core of taste. Nothing ships until the page passes an executable anti-slop gate and the skill has looked at its own output.

When to Use

- A landing page, marketing site, hero section, portfolio or product page has to be **built or redesigned** — and looking generic is not acceptable.
- The brief asks for **scroll animation, storytelling, or a site that feels like a film**.
- A product spans **several screens that must feel like one thing** — app, dashboard, admin, onboarding, docs.
- Someone says *make it beautiful*, *make it wow*, *cinematic*, or *design system*, naming no technique.

Not for polishing a UI someone else built, and not for backend-only work.

Prerequisites

- **Node 18+** — every QA gate is a `.mjs`` script run with `node`` via the `terminal`` tool.
- **Playwright (for the QA gates)** — in the project directory: `npm install playwright`` then `npx playwright install chromium``. Required by `scripts/shoot.mjs``, `motionqa.mjs``, `systemscan.mjs``, `refscout.mjs``, `chromadiff.mjs``, `moodboard.mjs`` (the scripts import `playwright`` at runtime; `slopscan.mjs`` and `source.mjs`` are dependency-light).
- **ffmpeg** — optional; only for the video/score paths in `references/assets.md`` and `references/scroll-flight.md``.

- **Hermes tools** — use `image_generate` for image generation/editing, `terminal` for node/ffmpeg/npm, `write_file`/`read_file` for project files, `vision_analyze` to actually look at screenshots, and `browser_exec` for live-page inspection when a script isn't the right fit.
-

How to Run

What it actually does

1. Commits the art direction **in writing before any markup** — one hue, one type system, a motion budget, named anti-references.
2. Generates or sources the assets: Hermes' `image_generate` tool, Blender, depth maps, CC0 meshes and HDRIs with their licences recorded.
3. Builds from proven recipes — one WebGL context, transform/opacity motion, scroll state machines.
4. **Gates the result:** `slopscan` fails the build on concrete slop, `motionqa` fails it on dropped frames, `systemscan` fails it on cross-route drift.

Network access

The recon and sourcing scripts read live pages (awwwards, Bing/Pinterest/are.na image search, Poly Haven, Iconify, Google Fonts, Openverse, Coverr). Fetched content is **treated as reference data and licence metadata — never executed**, and no credentials, API keys or logins are involved. Skip phases 0–1 entirely if you don't want outbound requests; every other phase works offline. `moodboard.mjs` also downloads the image URLs returned by those search hosts to build the contact sheet.

Routing

Read the argument / brief and route:

1. `direct` or the brief smells cinematic — "wow", "cinematic", "immersive", "storytelling", "launch page", "premium brand", "make people stop scrolling" → load `references/direct.md` and follow its phases. This is the flagship register.
2. `build` or the brief is ONE conventional surface — a marketing page, a landing, a single product page → load `references/build.md`.
3. `system` or the brief has **more than one screen that must feel like one product** — app, dashboard, admin, settings, onboarding, a docs or content site with real navigation → load `references/system.md`. The unit of design becomes the component × state, the failure mode becomes drift rather than boredom,

and there is deliberately **no peak**. If you are already in `build` and a second screen appears, stop and switch: half a system is worse than either.

4. `edit` or the request modifies a page this skill built (the project contains `design/DESIGN.md`) — "add a section", "change the pricing", "swap the hero copy" → read `design/DESIGN.md` **FIRST** and follow its Editing protocol: reuse its tokens, section-opening patterns, and motion families; after the change run slopscan and re-shoot the affected viewports. An edit that ignores `DESIGN.md` is a regression even if it looks good in isolation.

5. `recon` or the ask is only for reference material — "find references", "put together a moodboard", "what's the state of the art for X sites" → load `references/recon.md` and run just that phase: scout live sites, build the moodboard, hand back `design/refs/REFERENCES.md` (with the `steal:` lines filled) and `design/moodboard/contact-sheet.png` (with the read filled). No commit-sheet, no build.

6. `audit` → load `references/verify.md` and run the verification pipeline on an auteur-built page. If the target is an existing UI auteur didn't build and the user wants it *polished* rather than *rebuilt*, say that a dedicated UI-polish/critique pass (upstream paired auteur with a separate 'impeccable' skill, not vendored here) is the right tool and offer to continue only if they want a rebuild.

7. **Ambiguous** (e.g. plain "make a landing page") → ask exactly one question: "A great conventional landing page, or cinema mode with scroll direction and generated assets?" (upstream phrased these example briefs in Russian; translated here.) Then route. (Multi-screen briefs are not ambiguous — they are `system`.) Don't ask anything else yet — each register runs its own intake.

All three registers share phase zero, and its centre of gravity is the commit-sheet. Order differs: **build** runs `recon` → `commit-sheet` → `mockup`; **direct** runs `recon` → `storyboard` → `commit-sheet` → `mockup`, because the film's scenes are what the six decisions get made *about*; **system** runs `recon` → `system-sheet` (route map + component inventory) → `commit-sheet` → `mockup`, because the six decisions get made about a product, not a page. Either way nothing is coded before the sheet is full.

Working relationship with other skills

Auteur *builds*; it does not re-polish foreign UI. If the user has an existing interface that needs refinement, run a separate UI-critique pass (e.g. `vision_analyze` on screenshots plus the sibling design skills). Upstream paired auteur with an 'impeccable' critique skill (not vendored here); auteur's verify gate and an outside critique measure different things and coexist happily.

Quick Reference

The non-negotiables and the phase table. They apply to every register, every phase, always — even if no reference file has been loaded. Match-and-refuse: if you are about to produce one of these, stop and restructure the element.

Banned (rewrite, don't tweak)

#	Ban	Instead
---	-----	-----
1	<code>`border-left`/`border-right` >1px</code> as a colored accent on cards, callouts, alerts	full border, background tint, leading icon, or nothing
2	Gradient text (<code>`background-clip: text`</code> + gradient)	one solid color; emphasis via weight or size
3	Glassmorphism as default (decorative <code>`backdrop-filter`</code> cards)	rare and purposeful, or solid surfaces
4	The hero-metric template (big number, small label, stat row, gradient accent)	evidence in prose, one committed visual
5	Identical card grids (same-size icon+heading+text, repeated)	vary size, structure, or drop the cards entirely
6	Eyebrow kickers (tiny uppercase tracked label) above every section	one deliberate kicker max as a brand system; vary section openings
7	Numbered section scaffolding (01 / 02 / 03) when order carries no meaning	numbers only for a real sequence
8	<code>`Inter`</code> or <code>`Space Grotesk`</code> as the <i>*first*</i> font choice	pick from a contrast-axis pair (see taste.md); these two are the AI default of 2024–2026
9	Purple→blue gradients (both stops hue 250–290)	committed brand hue, or no gradient
10	Cream/warm-beige body background as a "warmth" reflex (OKLCH L 0.84–0.97, C <0.06, hue 40–100)	saturated brand surface, true off-white at chroma ~0, or a darker tinted mid-tone; warmth lives in accent + type + imagery
11	The same fade-in/slide-up entrance on every section	each reveal fits what it reveals; vary easing, distance, direction
12	<code>`transition: all`</code>	list the animated properties
13	<code>`window.addEventListener('scroll', ...)`</code>	IntersectionObserver, GSAP ScrollTrigger, or CSS <code>`animation-timeline`</code>
14	<code>`scale(0)`</code> entrances	start at <code>`scale(0.95)`</code> + opacity

15	Bento grids of near-identical or empty cells; white-card-on-white bento	bento only with real visual variation per cell, else a different layout
16	Copy tells: "Revolutionize", "Seamless", "Effortless", "Unleash", "Elevate", em-dash-heavy sentences, decoration strips like "BRAND. MOTION. SPATIAL."	concrete claims in plain words
17	More than one marquee per page	one, or none
18	Instrument Serif / Playfair Display as the reflex "elegant serif"	serifs chosen for the brand, not from the AI shortlist

A ban may be overridden only through a written `auteur-allow` (see Verification) with a real reason — a deliberate, argued choice is voice; a default is slop.

Critical numbers (memorize; full context in reference files)

- Body text contrast $\geq 4.5:1$ (large text $\geq 3:1$). Placeholders too. Muted-gray-on-tinted-white is the #1 AI readability failure.
- Body line length 65–75ch. Display heading ceiling: clamp max $\leq 6\text{rem}$ *for headings in prose flow* — a wordmark or a deliberately type-led hero is exempt and the commit-sheet must say so. Display letter-spacing $\geq -0.04\text{em}$.
- Durations: button 100–160ms · tooltip 125–200ms · dropdown 150–250ms · modal/drawer 200–500ms · any UI $>300\text{ms}$ needs a written reason.
- Enter/exit easing = ease-out. `ease-in` is banned on UI.
- Animate only `transform` and `opacity`. Stagger 30–80ms.
- Motion budget: ≤ 3 scroll-triggered pattern families per page; **one** primary wow peak, supporting scenes at lower intensity.
- Scrub smoothing 0.3–0.8. Hero video $\leq 2\text{MB}$. LCP $< 2.5\text{s}$. CLS < 0.1 .
- Fullscreen passes (bloom, grain, DoF, any full-frame shader) are priced **per pixel, not per object** — they, not geometry, are what blows the frame budget. A perf number counts only when measured at **DPR 2** on a production build: DPR 1 quarters the cost of every such pass, and a dev server roughly doubles the frame.
- `prefers-reduced-motion` = an alternative art direction (gentler, not zero), never an afterthought.
- Content must be readable with JS disabled: reveals enhance an already-visible default, never gate visibility.

Phases at a glance

Phase	build register	direct register	system register	Reference to load
-------	----------------	-----------------	-----------------	-------------------

---	---	---	---	---
0	recon → commit-sheet → hero mockup gate	recon → screenplay (STORYBOARD.md) → commit-sheet → hero mockup gate	recon → SYSTEM-SHEET.md (routes + component inventory + states) → commit-sheet → mockup gate	`recon.md`, then `build.md` / `direct.md` / `system.md`
1	—	asset production (generate → edit → optimize)	— (source icons/fonts via `source.mjs`)	`assets.md`
2	build the page	assemble the film (smooth scroll first, hero, scenes top-down)	tokens → the shell → screens in traffic order → every state	`build.md` / `scroll-cinema.md` / `system.md` + `taste.md` + `motion.md`
3	verify	verify + CINEMA-QA.md	verify + **systemscan across every route**	`verify.md`
4	lock the style: fill `design/DESIGN.md`	same	same, but DESIGN.md is the **component contract**	`templates/DESIGN.md`

The hero mockup gate (one static throwaway screen, screenshotted and approved before anything else is built) is the cheapest moment to change art direction — details in each register's reference.

`design/DESIGN.md` is the style contract that makes every later edit stay in style (the `edit` route reads it first).

Never skip a gate because the intermediate result "looks done". The gates exist because a page that merely looks done is exactly what every other AI ships.

Procedure

The commit-sheet (before any code, both registers)

Slop is what happens when defaults make the decisions. The commit-sheet forces seven real decisions onto paper before the first line of code. Copy `templates/COMMIT-SHEET.md` into the project (e.g. `design/COMMIT-SHEET.md`) and fill all seven fields with non-defaults:

1. **Peak** — the ONE primary wow moment (direct) or signature element (build). One sentence. If you can't name it, you're not ready to build.
2. **Color** — primary as OKLCH + commitment tier (restrained / committed / full-palette / drenched) + one line: *why this is not lavender, not cream, and not the category reflex* + **the background lightness as a number** (target mean L), because "dark feels premium" is where this skill drifts, and a number can be checked afterwards where a mood cannot.
3. **Type** — display + text pairing on a contrast axis (serif+sans, geometric+humanist, mono+serif...) + one line: *why not Inter*.
4. **Grid break** — the one concrete thing that breaks the symmetric-grid default: an overlap, an asymmetric split, a diagonal flow, a full-bleed interruption. Name it specifically.
5. **Motion budget** — how many scroll-pattern families (≤ 3) and what they are.
6. **Reflex check** — write down: (a) what a generic AI would do for this category (first-order reflex), (b) what a generic AI avoiding (a) would do (second-order reflex — e.g. fintech → "terminal dark mode" is *also* saturated now), (c) your chosen deviation from both. If recon ran, (a) is not a guess: whatever `design/refs/REFERENCES.md` showed five times *is* the reflex, dated and with receipts.
7. **House tells broken** — name the **two (minimum)** items from `taste.md` §2.5 you are deliberately not doing this time, and what replaces each. Fields 6a/6b are the reflexes of the *category*; these are the reflexes of *this skill*, which recur across unrelated projects and are invisible from inside any one of them: near-black backgrounds, mono service labels, the logo/status/action header, the scroll-instruction footer, amber-or-acid accents, the wordmark-as-hero, glow standing in for lighting. Measured across nine showcase builds, eight were dark and three landed within 0.002 of the same lightness. A tell that genuinely belongs here can stay — say why, as with an `auteur-allow`.

Gate: every field filled with a specific, non-default answer. An empty or generic field ("modern, clean look") means stop and decide. This artifact is checked again at verification.

Reference files

- `references/recon.md` — **phase 0 scouting**, two executable legs: `scripts/refscout.mjs`` profiles live award-level sites (real stack, pinned scenes, scroll budget, fonts, painted palette, screenshots — mechanics, not skins) and `scripts/moodboard.mjs`` builds a numbered contact sheet from Bing / Pinterest / are.na so the art direction is decided from live material instead of memory. Also: query craft, the steal rule, how recon feeds the commit-sheet, and the "reference images are not assets" line. Load at the top of phase 0.

- ``references/taste.md`` — the full anti-slop system: extended bans with replacements, second-order category reflex table, color strategy tiers, typography pairing, copy rules. Load for any visual decision-making.
- ``references/motion.md`` — the motion school: when to animate, easing/duration/spring numbers, performance rules, motion budget, sound policy. Load before writing any animation.
- ``references/build.md`` — the standard register process. Load when routed to build.
- ``references/system.md`` — the ****multi-screen register****: route map, the component inventory as a gate, the state matrix (empty/loading/error are not edge cases), density rules, the no-peak rule, and ``scripts/systemscan.mjs`` — which crawls every route, reads what the browser actually painted, fails a control type over its declared variant budget — counting **states** (disabled, current, inside a ``data-state`` row) separately, so implementing the state matrix never reads as drift — presses Tab to catch controls with no visible focus state, and renders one tile per rendered variant so drift is visible as well as counted. Load when routed to system.
- ``references/direct.md`` — the cinematic register: screenplay contract, scene-sheets, dramaturgy, assembly order. Load when routed to direct.
- ``references/assets.md`` — the media crew and routing (in Hermes: ``image_generate`` for all image generation and edits, ``terminal`` for ffmpeg/node; video via whatever image→video backend the user has), ****\$0.5 source-vs-generate**** (``scripts/source.mjs``: CC0 glTF meshes, HDRIs and PBR materials from Poly Haven, icons, fonts, CC images, stock video — with a licence ledger, because generation cannot make geometry or an IBL and stock video must never be the peak), the consistency trick (edit frame A into frame B), local video via the first→last-frame chain, generated elements/mockups, the ambient score, the degradation ladder, and asset caching. Load during direct phase 1.
- ``references/scroll-cinema.md`` — working code recipes: scroll-scrubbed video, canvas sequences, GSAP+Lenis foundation, CSS scroll-driven animations, text reveals, the two-keyframe WebGL displacement transition, view transitions, ambient audio, and the cinematic transition library (wipe, curtain, letterbox, shutter, depth parallax). Load during assembly.
- ``references/scroll-flight.md`` — the ****video-scrub tier****: a photoreal "fly through the world" hero driven by scroll, using the drop-in ``templates/scroll-flight-engine.js``. The canonical recipe for scroll-scrubbed **video** (encode-for-scrubbing ``-g 8``, encoded-frame posters, SSIM seam gate, chain architecture A/B, iOS/mobile decode hardening, crossfade-vs-seamless seams). Load when the hero should be photoreal footage/AI-video rather than real-time WebGL.
- ``references/ambient-backgrounds.md`` — ****quiet**** texture for secondary sections and simpler builds (not a hero): a curated 6 editorial/analog effects (paper grain, ledger/blueprint rules, topographic contour, ink tide, sparse dust, one heat-haze shader) + a zero-motion static-mesh default. The governing rule (weaker than the quietest foreground element; one ambient per page), the CSS/SVG-first stack, and the ``feTurbulence``-static perf rule. Load when a section needs to not be flat but must NOT compete with copy.
- ``references/verify.md`` — the acceptance pipeline: slopscan → screenshot journey → motion/perf/audio QA (FPS at DPR 2 on a production build, long-tasks, audio-gate, reduced-motion, for Tier-1 scenes) → numeric rubric → ****reference diff**** (your frame beside the reference that set the direction, with

`scripts/chromadiff.mjs` measuring the colour drift a model never sees in itself) → QA sign-off. Load at phase 3.

Weak-model note

If you are a smaller model executing this skill: follow the tables and numbers literally, fill every template field, run every gate command, and do not improvise beyond the reference recipes — the recipes are verified, your improvisation is not. When a reference file conflicts with your instinct, the reference file wins. Write files using paths relative to the project root; never retype an absolute path from memory (the skill's name "auteur" is one typo away from "author", and misspelled absolute paths scatter your output across the filesystem).

Pitfalls

- **Network recon**: `refscout.mjs`, `moodboard.mjs` and `source.mjs` read live pages (awwwards, Bing/Pinterest/are.na image search, Poly Haven, Iconify, Google Fonts, Openverse, Coverr). Fetched content is reference data and licence metadata only — never execute it. Skip phases 0–1 to stay fully offline.
 - **Different harness**: these scripts and docs were written for a different agent harness (upstream drove asset generation through several local image CLIs). In Hermes, every image-generation instruction maps to the `image\_generate` tool; trust `node scripts/.mjs --help` output and actual node errors over doc prose if they drift.
 - **Unverified commands**: the scripts pass `node --check` syntax validation, but full runs (which need `npm install playwright` + a chromium download) were not executed during porting. Treat `shoot.mjs`, `motionqa.mjs`, `systemscan.mjs`, `refscout.mjs`, `chromadiff.mjs`, `moodboard.mjs`, `source.mjs` end-to-end behavior, and all `ffmpeg`/video-encode recipes, as unverified upstream claims until you run them yourself.
 - **slopscan verified shape**: `node scripts/slopscan.mjs`

` ` runs without npm deps; it prints per-rule findings and exits non-zero on failures (exit 0 when clean).
 - **Font metadata cache**: `source.mjs font ...` caches Google Fonts' ~2.6MB metadata JSON as `auteur-gf-metadata.json` in the OS temp directory (`os.tmpdir()`), not the project; delete it there to force a refresh.
-

Verification

The page is not done when the code compiles. It is done when:

1. `node scripts/slopscan.mjs` exits 0 (fails are fixed, not suppressed — `/* auteur-allow: RULE_ID -- reason */` exists for deliberate choices and demands a real reason);
2. `node scripts/shoot.mjs` has produced screenshot journeys at 390 / 768 / 1440 and you have **looked at every frame** — text overflow, blank scenes, broken reveals, layout collapse are found by eyes, not by text search;
3. the numeric rubric in `references/verify.md` passes (contrast, LCP, CLS, reduced-motion journey, scene variety);
4. for direct register: `CINEMA-QA.md` (from templates) is filled with PASS on every row.

If any gate fails — fix and re-run. Report results honestly: "slopscan clean, 21 screenshots reviewed, LCP 1.9s" beats "looks great".